

Ming Gong, Ph.D.

+1 206-833-0010 · Bellevue, WA · yiming10.13@163.com · Google Scholar

SUMMARY

Seasoned **applied-sciences and full-stack engineering leader** with a proven track record of building consumer-scale **search and AI assistant** products. Currently Head of ML Engineering at Atlassian, serving as the **single point of accountability for an extended org of 100+** (60+ direct reports plus dotted-line partner teams) across Jira Search, Knowledge Fundamentals, and Generative Answers & AI Search Experiences. Previously Group Senior Principal Applied Sciences Manager at Microsoft, built and led a 40+ team that took Bing Question-Answering, Knowledge Content and Experiences, and Copilot Search from English-only to **100+ languages and 200+ markets**, and shipped **Copilot Search as a Bing AI vertical one month ahead of Google's AI Mode**.

End-to-end search and AI expertise spanning query understanding, search relevance & ranking modeling, extractive and generative answers, agentic search frameworks & experiences, personalization, LLM fine-tuning & post-training, and large-scale ML serving. Uniquely multi-dimensional: **applied sciences & full-stack engineering, AI products & research** (50+ publications, 10K+ citations, 20+ patents), **consumer & enterprise products**.

SELECTED HIGHLIGHTS

- **Consumer-scale impact:** Knowledge Card on Bing — 150M entities, 80% query traffic, **32M DAU**, +3.0% WPSBS; certified as a Bing "WoW" feature. Generative Answers — 20% query coverage globally, +2.0% WPSBS.
- **Generative AI assistant:** Drove Copilot Search (multi-turn AI search vertical on Bing), Generative SERP, and Generative Answers. Established LLM post-training and SLM fine-tuning framework; **<5s first-token latency** at production scale.
- **Global & multilingual:** Pioneered cross-lingual scaling methodology for QnA across 100+ languages, 200+ markets — **2x Microsoft MLADS Distinguished Contribution Awards** (first-ever back-to-back recipient; CTO Office global awards).
- **Ship velocity:** Jira Search MVP shipped **4 months** after joining Atlassian, full version launch to all customers within 12 months, driving **3x Search MAU**; Generative Answers in search and AI Definitions drove AI MAU from 100K to **1.2M**; Enterprise Generative AI Search experiences POC in **<1 month**, MVP GA in early 2026.
- **Org scale:** Single POC for an **extended org of 100+** at Atlassian (60+ direct + dotted-line partner teams); previously built a 40+ team at Microsoft from scratch across two sites.
- **Talent development:** Founded and led **Microsoft AI School APRD** for 8 years, leading a **50+ operation v-team** — 10K+ participants, 1K+ honored graduates; the most influential AI training program in Microsoft Asia-Pacific R&D.
- **Research recognition:** Best Paper Award, AI4Research at IJCAI '24; 50+ publications, 10K+ citations.

EXPERIENCE

Atlassian (US)

Feb 2025 – Present

Head of Machine Learning Engineering — Knowledge AI

Bellevue, WA

- Single point of accountability for the Knowledge AI charter, leading an **extended org of 100+** — **60+ direct reports** (SWEs and MLEs, full-stack and globally distributed) plus dotted-line partner teams across Jira Search, Knowledge Fundamentals, and Generative Answers & AI Search Experiences. *[Ownership / Hire and Develop]*
- **Shipped Jira Search MVP within 4 months** of joining, full version launch to all customers within 12 months; rebuilt ranking stack and search UX, driving **3x growth in Search MAU**; premiered at Team '25. *[Bias for Action / Deliver Results]*
- Generative Answers in search and AI Definitions drove sustained **AI MAU from 100K to 1.2M** via novel generative UI framework. *[Customer Obsession / Invent and Simplify]*
- Built core **knowledge fundamentals** (collaboration graph, entity linking, entity recommendation) now adopted by multiple product teams for personalized features. *[Think Big / Ownership]*
- Launched **Jira Agents in RoVo Chat** within 3 months, becoming a primary driver of RoVo adoption. *[Customer Obsession]*
- Initiated and led **Enterprise Generative AI Search** from zero to GA — POC in <1 month, MVP shipped in early 2026. The novel **Generative UI layer** developed for this product was adopted as Atlassian's standard generative presentation

framework, now being integrated across all GenAI products and features. [\[Invent and Simplify / Think Big\]](#)

Microsoft — Search Technology Center Asia (Bing, Edge, Copilot)

Aug 2013 – Dec 2024

Group Senior Principal Applied Sciences Manager (2019–2024) · Senior Dev Manager (2018–2019) · Senior Tech Lead (2016–2018) · SWE II (2013–2015) Beijing, China

- **Built and led an org of 40+ engineers and applied scientists** across Beijing and Suzhou, spanning Bing Search, Copilot, Cross-Lingual QA, Knowledge Graph, Recommendation, and Edge browser AI. [\[Hire and Develop\]](#)
- **Copilot Search (Bing AI vertical, 2024–2025)**: delivered multi-turn AI search experience unifying GenQnA, GenSERP, and Magazine; built scalable LLM generation pipelines and task-specific SLMs (incl. R1-distilled). **Shipped April 2025 — one month before Google's AI Mode.** [\[Bias for Action / Think Big\]](#)
- **Generative SERP (2023–2024)**: reinvented Bing SERP as "AI Masterpieces" — automated deep-intent understanding and multi-source synthesis with hallucination control. Shipped MVP at 1.0% query coverage with +0.2 WPSBS; V2 SLM pipeline scaled to 5–10%. [\[Invent and Simplify\]](#)
- **Generative Answers on Bing SERP (2023–present)**: **20% query coverage on Bing global markets, +2.0% WPSBS** via streaming generative answers and click-through to Bing Chat. [\[Customer Obsession / Deliver Results\]](#)
- **Knowledge Card for 150M Entities (2021–2023)**: multi-modal entity aggregation via template + MRC DL extraction and LLM generation. **80% Bing query traffic, 32M DAU, +3.0 WPSBS**; certified as a Bing "WoW" feature. [\[Customer Obsession / Deliver Results\]](#)
- **Segment Entity Search & Recommendation (2021–present)**: shipped recipe, job, and movie answers — **beat Google on relevance and WPSBS** on recipe and job; closed WPSBS gap with Google on movie recommendation. [\[Are Right, A Lot\]](#)
- **Microsoft Copilot on Edge (2023)**: long-document summarization and chat over PDFs and long-form web pages — 1000+ page docs at **95%+ accuracy**; demoed at Microsoft Copilot 2023 event; shipped to Edge and Windows Copilot. [\[Deliver\]](#)
- **NLU Platform "Orca" (2020–2025)**: unified query understanding pipeline (domain, intent, slots, topics) powering answer triggering. **95%+ of Bing answers (coverage ≥ 0.05%)** onboarded; 60% onboarding answers closed gap with Google. [\[Invent & Simplify\]](#)
- **Globalization of Bing Q&A (2016–2022)**: led the team to scale QnA from English-only to **100+ languages and 200+ markets**; proposed systematic cross-lingual scaling methodology — **2× MLADS Distinguished Contribution Award** (first-ever back-to-back recipient; CTO Office global awards). Shipped 10+ DNN releases replacing legacy GBDT ranker (400+ hand-crafted features), won WPSBS vs. Google. Authored novel model compression that powered the **first BERT model deployed in Bing** (within 1 month of BERT open-source) — +3.0% coverage, +5.0% precision. [\[Think Big / Are Right, A Lot\]](#)

PEOPLE & ORG DEVELOPMENT

- **Founder & Head, Microsoft AI School — Asia-Pacific R&D Group (2017–2024)**. Led a **50+ operation v-team**; drove planning, curriculum, and execution of all AI learning events for the region. **10K+ participants, 1K+ honored graduates** (with strict graduation criteria); record-breaking participation 8 years running. [\[Hire and Develop\]](#)
- **Invited instructor, LinkedIn Global Academy for Chinese Programmers (2022–2023)**: 4-class series on "Intelligent QA Systems for Search Engines" — **50M impressions, 98% satisfaction**.
- **Invited speaker, InfoQ AICon Global AI Conference (2021)**: represented Microsoft China on globalization of Bing QA.

EDUCATION

Ph.D., Graphics and Visual Computing — Institute of Computing Technology, Chinese Academy of Sciences 2008 – 2013

SELECTED AWARDS

- **Best Paper Award, IJCAI '24 Workshop** — "Step-Back Profiling: Distilling User History for Personalized Scientific Writing."
- **2021 Fall Microsoft MLADS Distinguished Contribution Award (2 / 250+)**. Broke conference history to receive the award twice consecutively.
- **2021 Spring Microsoft MLADS Distinguished Contribution Award (2 / 300+)**.
- **2023 Microsoft Global Hackathon** — Edge Theatre, Hack for Web Experiences: **Best Overall Hack**; Greater China Region (Beijing): Most Popular Hack.

- 2023 Microsoft Global Hackathon — WebXT China, DeKG (decentralized knowledge ecosystem): Best Innovativeness & Originality Hack; Beijing Outstanding Project Award.
- Excellent Operation Award of AI School China — 2018, 2019, 2020, 2021, 2022, 2023.
- Microsoft AI Core Q3 FY18 Greatness Award (2018).

SELECTED PUBLICATIONS (50+ papers, 10,000+ citations on Google Scholar — full list on [scholar profile](#))

► Cross-Lingual NLU & Multilingual Scaling

- Nuo Chen, Linjun Shou, **Ming Gong**, Jian Pei, Daxin Jiang. "Bridging the Gap between Language Models and Cross-Lingual Sequence Labeling." *NAACL* 2022.
- Nuo Chen, Linjun Shou, **Ming Gong**, Jian Pei. "From Good to Best: Two-Stage Training for Cross-lingual Machine Reading Comprehension." *AAAI* 2022.
- Shining Liang, **Ming Gong**, Jian Pei, Linjun Shou, Daxin Jiang. "Reinforced Iterative Knowledge Distillation for Cross-Lingual Named Entity Recognition." *KDD* 2021.
- Yingmei Guo, Linjun Shou, Jian Pei, **Ming Gong**, et al. "Learning from Multiple Noisy Augmented Data Sets for Better Cross-Lingual Spoken Language Understanding." *EMNLP* 2021.
- Shining Liang, Linjun Shou, Jian Pei, **Ming Gong**, et al. "CalibreNet: Calibration Networks for Multilingual Sequence Labeling." *WSDM* 2021.
- Nuo Chen, Linjun Shou, Tengtao Song, **Ming Gong**, et al. "Structural Contrastive Pretraining for Cross-Lingual Comprehension." *ACL* 2023.
- Nuo Chen, Zinan Zheng, Ning Wu, **Ming Gong**, et al. "Breaking Language Barriers in Multilingual Mathematical Reasoning." *arXiv* 2023.
- Junhao Liu, Linjun Shou, Jian Pei, **Ming Gong**, et al. "Cross-lingual Machine Reading Comprehension with Language Branch Knowledge Distillation." *COLING* 2020.
- F. Yuan, Linjun Shou, X. Bai, **Ming Gong**, et al. "Enhancing Answer Boundary Detection for Multilingual Machine Reading Comprehension." *ACL* 2020.
- H. Huang, Yaobo Liang, N. Duan, **Ming Gong**, et al. "Unicoder: A Universal Language Encoder by Pre-training with Multiple Cross-lingual Tasks." *EMNLP/IJCNLP* 2019.
- Linjun Shou, **Ming Gong**, Jian Pei, et al. "Language Scaling: Applications, Challenges and Approaches." *KDD Tutorial* 2021.
- Linjun Shou, **Ming Gong**, Jian Pei, et al. "Scaling out NLP Applications to 100+ Languages." *WWW Tutorial* 2021.

► Search Relevance, Retrieval & Question Answering

- Shengyao Zhuang, Linjun Shou, Jian Pei, **Ming Gong**, et al. "Typos-aware Bottlenecked Pre-Training for Robust Dense Retrieval." *SIGIR* 2023.
- Houxing Ren, Linjun Shou, Jian Pei, Ning Wu, **Ming Gong**, Daxin Jiang. "Lexicon-Enhanced Self-Supervised Training for Multilingual Dense Retrieval." *EMNLP* 2023.
- Houxing Ren, Linjun Shou, Ning Wu, **Ming Gong**, Daxin Jiang. "Empowering Dual-Encoder with Query Generator for Cross-Lingual Dense Retrieval." *EMNLP* 2023.
- Shunyu Zhang, Yaobo Liang, **Ming Gong**, Daxin Jiang, Nan Duan. "Modeling Sequential Sentence Relation to Improve Cross-lingual Dense Retrieval." *ICLR* 2023.
- Shunyu Zhang, Yaobo Liang, **Ming Gong**, Daxin Jiang, Nan Duan. "Multi-View Document Representation Learning for Open-Domain Dense Retrieval." *ACL* 2022.
- Junjie Huang, Wanjun Zhong, Qian Liu, **Ming Gong**, et al. "Mixed-modality Representation Learning and Pre-training for Joint Table-and-Text Retrieval in OpenQA." *Findings of EMNLP* 2022.
- Linjun Shou, S. Bo, Fei-Xiang Cheng, **Ming Gong**, et al. "Mining Implicit Relevance Feedback from User Behavior for Web Question Answering." *KDD* 2020.
- Xingyao Zhang, Linjun Shou, Jian Pei, **Ming Gong**, et al. "A Graph Representation of Semi-structured Data for Web Question Answering." *COLING* 2020.
- Xuguang Wang, Linjun Shou, **Ming Gong**, Nan Duan, Daxin Jiang. "No Answer is Better Than Wrong Answer: A Reflection Model for Document Level Machine Reading Comprehension." *EMNLP* 2020.
- Nuo Chen, Linjun Shou, **Ming Gong**, et al. "Alleviating Over-smoothing for Unsupervised Sentence Representation." *ACL* 2023.

► Knowledge Graphs, Entity Understanding & Information Extraction

- Zilin Xiao, **Ming Gong**, Jie Wu, et al. "Instructed Language Models with Retrievers Are Powerful Entity Linkers." *EMNLP* 2023.
- Zilin Xiao, Linjun Shou, Xingyao Zhang, Jie Wu, **Ming Gong**, et al. "Coherent Entity Disambiguation via Modeling Topic and Categorical Dependency." *Findings of EMNLP* 2023.
- Zimeng Li, Bo Shao, Linjun Shou, **Ming Gong**, et al. "WIERT: Web Information Extraction via Render Tree." *AAAI* 2023.

- Changzhi Sun, Yeyun Gong, Yuanbin Wu, **Ming Gong**, et al. "Joint Type Inference on Entities and Relations via Graph Convolutional Networks." *ACL* 2019.

► **Model Compression, Knowledge Distillation & Efficient Serving**

- Ze Yang, Linjun Shou, **Ming Gong**, Wutao Lin, Daxin Jiang. "Model Compression with Two-stage Multi-teacher Knowledge Distillation for Web Question Answering System." *WSDM* 2020.
- Fei Yuan, Linjun Shou, Jian Pei, Wutao Lin, **Ming Gong**, Daxin Jiang. "Reinforced Multi-Teacher Selection for Knowledge Distillation." *AAAI* 2021.

► **LLM, Agentic AI & Multimodal Pre-training**

- Nuo Chen, Ning Wu, Shining Liang, **Ming Gong**, et al. "Is Bigger and Deeper Always Better? Probing LLaMA Across Scales and Layers." *arXiv* 2023.
- Ning Wu, **Ming Gong**, Linjun Shou, et al. "Large Language Models are Diverse Role-Players for Summarization Evaluation." *NLPCC* 2023.
- Yaobo Liang, Chenfei Wu, Ting Song, et al., **Ming Gong**, Nan Duan. "TaskMatrix.AI: Completing Tasks by Connecting Foundation Models with Millions of APIs." *Intelligent Computing* 2023.
- Gen Li, Nan Duan, Yuejian Fang, **Ming Gong**, Daxin Jiang. "Unicoder-VL: A Universal Encoder for Vision and Language by Cross-modal Pre-training." *AAAI* 2019.
- **Ming Gong**, Linjun Shou, Wutao Lin, et al. "NeuronBlocks — Building Your NLP DNN Models Like Playing Lego." *EMNLP/IJCNLP* 2019.

► **Recommendation & Personalization**

- Ning Wu, **Ming Gong**, Linjun Shou, Jian Pei, Daxin Jiang. "RUEL: Retrieval-Augmented User Representation with Edge Browser Logs for Sequential Recommendation." *CIKM* 2023.
- Hanshuang Tong, Jun Li, Ning Wu, **Ming Gong**, et al. "Plutos: Towards Interpretable Stock Movement Prediction with Financial Large Language Model." *arXiv* 2024.
- Junjie Huang, Duyu Tang, Linjun Shou, **Ming Gong**, et al. "CoSQA: 20,000+ Web Queries for Code Search and Question Answering." *ACL* 2021.
- Junwei Liao, Yu Shi, **Ming Gong**, et al. "Generating Human Readable Transcript for Automatic Speech Recognition with Pre-trained Language Model." *ICASSP* 2021.

SELECTED PATENTS (20+ granted; selection below)

- Personalized Multilingual Content Recommendation based on Robust Feature Network. 2023. (MS#412659-CN-NP)
- Lexicon-Enhanced Self-Supervised Training for Multilingual Dense Retrieval. 2022. (MS#412496-CN-NP)
- Sequential Recommendation based on Cross-Domain Behavior Data. 2023.
- Temporal Co-Contrastive Learning-Based Node Representation Generation. 2022.
- Sequential Recommendation Based on Dual Contrastive Interest Learning. 2022.
- Sentence Representation Generation for Cross-lingual Retrieval. 2022.
- Content Recommendation Based on Graph Enhanced Collaborative Filtering. 2022.
- Controversial Poll Auto Generation by Generation Model and Crowdsourcing. 2021.
- Knowledge-enhanced Content-Aware Collaborative Filtering for Recommendation. 2021.
- From Good to Best: Two-Stage Training for Cross-lingual Machine Reading Comprehension. 2021.
- Graph based Fact Clustering for Diversity and Freshness. 2021.
- Representation Learning of Cross-Language Texts. 2021. (MS#409565-CN-NP)
- Hierarchical Representation Learning of User Interest. 2021. (MS#410462-CN-NP)
- Question Generation from Queries. 2021. (MS#410241-CN-NP)
- Multi-Model Joint Denoising Training. 2021. (MS#409566-CN-NP)
- Representation Learning of Semi-structured Data. 2020. (MS#408908-CN-NP)
- Model Selection Learning for Knowledge Distillation. 2020. (MS#408426-CN-NP)
- Providing QA Training Data and Training a QA Model based on Implicit Relevance Feedbacks. 2020. (MS#407927-CN-NP)
- Contrastive Learning for Question Answering. 2020. (MS#407988-CN-NP)
- Model Generation based Model Compression. 2019. (MS#406805-CN-NP)
- Cross-Lingual Task Training. 2019. (MS#406497-CN-NP)
- Assertion-based Question Answering with Question-Aware Open Information Extraction. 2019.